

Lecture 13 : Neural Networks

- We already looked at two basic parametric models
 - linear regression
 - logistic regression
- A **neural network**, in some sense, extends these basic models by stacking multiple copies of these models to construct a **hierarchical model**
- This hierarchical model can describe **more complicated relationships** between inputs and outputs than a linear or logistic regression
- **Deep learning** is a subfield of machine learning that deals with such hierarchical machine learning models

Neural Network Model

- Earlier we introduced the concept of **non-linear parametric functions** for modelling the relationship between input variables x_1, \dots, x_p and output y

$$\hat{y} = f_{\underline{\theta}}(\underline{x}) = f_{\underline{\theta}}(x_1, x_2, \dots, x_p)$$

The function f is "parameterized" by $\underline{\theta}$

- Such a non-linear function $f_{\underline{\theta}}(\cdot)$ can be created in many ways
- In neural network, the strategy is to use several **layers** of linear regression models and non-linear **activation functions**

- In linear regression, we wrote the model prediction as:

$$\hat{y} = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_p x_p$$


$$\underline{x} = \begin{bmatrix} 1 \\ x_1 \\ \vdots \\ x_p \end{bmatrix} \quad \text{input vector}$$

- We start the description of a neural network model with linear regression model

$$\hat{y} = \overset{\theta_0}{b} + \overset{\theta_1}{w_1} x_1 + \overset{\theta_2}{w_2} x_2 + \dots + \overset{\theta_p}{w_p} x_p$$

w_1, w_2, \dots, w_p are called the **weights** and b is the **offset**
 we use this notation because it is more popular in neural networks


- Graphical representation



$\xrightarrow{w_j}$ - Each link is associated with a parameter

\hat{y} - Output \hat{y} is described as sum of all terms

$$\hat{y} = w_1 x_1 + \dots + w_p x_p + b$$



$$\hat{y} = w_1 x_1 + \dots + w_p x_p + b$$

Describes a linear relationship

- How to describe a non-linear relationship between

$$\underline{x} = \begin{bmatrix} 1 \\ x_1 \\ \vdots \\ x_p \end{bmatrix} \quad \text{and} \quad \hat{y} ?$$


- Use activation function $h: \mathbb{R} \rightarrow \mathbb{R}$

$$\hat{y} = \sigma(w_1 x_1 + \dots + w_p x_p + b)$$

This now becomes a generalized linear model


- Common choices of activation functions are:

Logistic : $\sigma(z) = \frac{1}{1 + e^{-z}}$
(Sigmoid)



Logistic: $\sigma(z) = \frac{1}{1 + e^{-z}}$

ReLU: $\sigma(z) = \max(0, z)$
(standard choice)



ReLU: $\sigma(z) = \max(0, z)$

$$\hat{y} = \sigma(w_1 x_1 + \dots + w_p x_p + b)$$

This now becomes a generalized linear model

- This generalized linear model is very simple
- Cannot describe very complicated relationships between input \underline{x} and output \hat{y}

- How can we extend this simple model to increase the flexibility?

Stack several generalized linear models in a sequential construction

Ex:

One-layer
of stacking
(hidden layer)



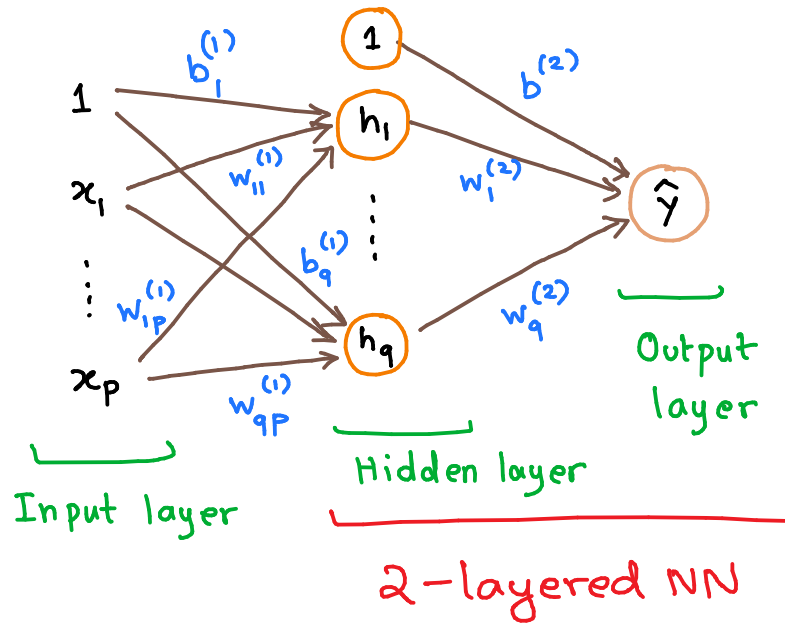
$$\begin{cases} h_1 = \sigma(w_{11}^{(1)} x_1 + w_{12}^{(1)} x_2 + \dots + w_{1p}^{(1)} x_p + b_1^{(1)}) \\ h_2 = \sigma(w_{21}^{(1)} x_1 + w_{22}^{(1)} x_2 + \dots + w_{2p}^{(1)} x_p + b_2^{(1)}) \\ \vdots \\ h_q = \sigma(w_{q1}^{(1)} x_1 + w_{q2}^{(1)} x_2 + \dots + w_{qp}^{(1)} x_p + b_q^{(1)}) \end{cases}$$

Output layer



$$\hat{y} = w_1^{(2)} h_1 + w_2^{(2)} h_2 + \dots + w_q^{(2)} h_q + b^{(2)}$$

Two-layer Neural Network

- Each link (arrow) is associated with a weight



- Each circular node is associated with its own bias term b



$$h_1 = \sigma(w_{11}^{(1)} x_1 + w_{12}^{(1)} x_2 + \dots + w_{1p}^{(1)} x_p + b_1^{(1)})$$

$$h_2 = \sigma(w_{21}^{(1)} x_1 + w_{22}^{(1)} x_2 + \dots + w_{2p}^{(1)} x_p + b_2^{(1)})$$

\vdots

$$h_q = \sigma(w_{q1}^{(1)} x_1 + w_{q2}^{(1)} x_2 + \dots + w_{qp}^{(1)} x_p + b_q^{(1)})$$

$$\hat{y} = w_1^{(2)} h_1 + w_2^{(2)} h_2 + \dots + w_q^{(2)} h_q + b^{(2)}$$

Vectorized Representation

- The 2-layered neural network can be more compactly written using matrix notation:

$$\underline{\underline{W}}^{(1)} = \begin{bmatrix} w_{11}^{(1)} & \dots & w_{1p}^{(1)} \\ \vdots & & \vdots \\ w_{q1}^{(1)} & \dots & w_{qp}^{(1)} \end{bmatrix}_{q \times p}, \quad \underline{b}^{(1)} = \begin{bmatrix} b_1^{(1)} \\ \vdots \\ b_q^{(1)} \end{bmatrix}_{q \times 1}, \quad \underline{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_p \end{bmatrix}_{p \times 1}$$

$$\underline{\underline{W}}^{(2)} = \begin{bmatrix} w_1^{(2)} & \dots & w_q^{(2)} \end{bmatrix}_{1 \times q}, \quad \underline{b}^{(2)} = \begin{bmatrix} b^{(2)} \end{bmatrix}_{1 \times 1}, \quad \underline{h} = \begin{bmatrix} h_1 \\ \vdots \\ h_q \end{bmatrix}_{q \times 1}$$

Compact
representation

$$\underline{h}_{q \times 1} = \sigma \left(\underbrace{\underline{\underline{W}}^{(1)} \underline{x}}_{q \times 1} + \underline{b}^{(1)} \right)$$

$$\hat{y}_{1 \times 1} = \underline{\underline{W}}^{(2)}_{1 \times q} \underline{h}_{q \times 1} + \underline{b}^{(2)}_{1 \times 1}$$


$$\underline{\Theta} = \begin{bmatrix} \text{vec}(\underline{\underline{W}}^{(1)}) \\ \underline{b}^{(1)} \\ \text{vec}(\underline{\underline{W}}^{(2)}) \\ b^{(2)} \end{bmatrix}$$

$$\hat{y} = f_{\underline{\Theta}}(\underline{x})$$


Note: The activation function $\sigma(\cdot)$ operates element-wise

Deep neural network


- The 2-layered neural network is called **shallow NN** (because it has one hidden layer)
- The real flexibility of neural network comes when we have more than one hidden layer.
- Stacking multiple hidden layers leads to a **DEEP** neural network



...



Deep Neural Network (Feed-forward Neural Network — FNN)



Mathematical
representation
of FNN:

$$\underline{h}^{(1)} = \sigma(\underline{\underline{W}}^{(1)} \underline{x} + \underline{b}^{(1)})$$

$$\underline{h}^{(2)} = \sigma(\underline{\underline{W}}^{(2)} \underline{h}^{(1)} + \underline{b}^{(2)})$$

\vdots

$$\underline{h}^{(L)} = \sigma(\underline{\underline{W}}^{(L-1)} \underline{h}^{(L-1)} + \underline{b}^{(L-1)})$$


$$\hat{y} = \underline{\underline{W}}^{(L)} \underline{h}^{(L-1)} + \underline{b}^{(L)}$$

Vectorization over datapoints

- During training, the neural network model is used to compute the predicted output for several input points $\{\underline{x}_i\}_{i=1}^N$

Superscript (l) \rightarrow will denote layer #

- The 2-layer neural network



$$\begin{aligned} \underline{h}_{q \times 1} &= \sigma \left(\underline{w}^{(1)}_{q \times p} \underline{x}_{p \times 1} + \underline{b}^{(1)}_{q \times 1} \right) \\ \hat{y}_{1 \times 1} &= \underline{w}^{(2)}_{1 \times q} \underline{h}_{q \times 1} + \underline{b}^{(2)}_{1 \times 1} \end{aligned}$$

(column vector)

$$\begin{aligned} \underline{h}_i^T &= \sigma \left(\underline{x}_i^T \underline{w}^{(1)T} + \underline{b}^{(1)T} \right) \\ \hat{y}_i &= \underline{h}_i^T \underline{w}^{(2)T} + \underline{b}^{(2)T} \end{aligned}$$


(row vector)

$i = 1, 2, \dots, n$

Transpose

Vectorization over datapoints

- During training, the neural network model is used to compute the predicted output for several input points $\{y_i, \underline{x}_i\}_{i=1}^N$
- The 2-layer neural network



$$\underline{h}_{q \times 1} = \sigma \left(\underline{W}^{(1)} \underline{x} + \underline{b}^{(1)} \right)$$


$q \times 1$ $q \times 1$ $q \times 1$

column vector

$$\hat{y}_{1 \times 1} = \underline{W}^{(2)} \underline{h} + \underline{b}^{(2)}$$

1×1 $1 \times q$ $q \times 1$ 1×1

\rightarrow



$$\underline{h}_i^T = \sigma \left(\underline{x}_i^T \underline{W}^{(1)T} + \underline{b}^{(1)T} \right)$$

$1 \times q$ $1 \times p$ $p \times q$ $1 \times q$

row vector

$$\hat{y}_i = \underline{h}_i^T \underline{W}^{(2)T} + \underline{b}^{(2)T}$$

1×1 $1 \times q$ $q \times 1$ 1×1

$i = 1, 2, \dots, n$

Transposed

- Similar to linear regression, we stack all data points in matrices

$$\underline{y} = \begin{bmatrix} y_1 \\ \vdots \\ y_N \end{bmatrix}_{N \times 1}, \quad \underline{X} = \begin{bmatrix} \underline{x}_1^T \\ \vdots \\ \underline{x}_N^T \end{bmatrix}_{N \times p}, \quad \hat{\underline{y}} = \begin{bmatrix} \hat{y}_1 \\ \vdots \\ \hat{y}_N \end{bmatrix}_{N \times 1}, \quad \underline{H} = \begin{bmatrix} \underline{h}_1^T \\ \vdots \\ \underline{h}_N^T \end{bmatrix}_{N \times q}$$

each row represents a training data point

$$\underline{H} = \sigma \left(\underline{X} \underline{W}^{(1)T} + \underline{b}^{(1)T} \right)$$

$N \times q$ $N \times p$ $p \times q$ $1 \times q$

$$\hat{\underline{y}} = \underline{H} \underline{W}^{(2)T} + \underline{b}^{(2)T}$$

$N \times 1$ $N \times q$ $q \times 1$ 1×1

added to each row

$$\underline{\underline{H}} = \sigma \left(\underline{\underline{X}} \underline{\underline{W}}^{(1)T} + \underline{\underline{b}}^{(1)T} \right)$$

$$\underline{\underline{\hat{Y}}} = \underline{\underline{H}} \underline{\underline{W}}^{(2)T} + \underline{\underline{b}}^{(2)T}$$



$$\underline{\underline{H}} = \sigma \left(\underline{\underline{X}} \tilde{\underline{\underline{W}}}^{(1)} + \tilde{\underline{\underline{b}}}^{(1)} \right)$$

$$\underline{\underline{\hat{Y}}} = \underline{\underline{H}} \tilde{\underline{\underline{W}}}^{(2)} + \tilde{\underline{\underline{b}}}^{(2)}$$

- These vectorized equations would be used in implementation in PyTorch, Tensorflow, etc.

- During programming, you may consider using the transposed versions of $\underline{\underline{W}}$ and $\underline{\underline{b}}$ as the weight matrix and bias vectors to avoid transposing them in each layer

$$\tilde{\underline{\underline{W}}}^{(1)} \leftarrow \underline{\underline{W}}^{(1)T}$$

$$\tilde{\underline{\underline{W}}}^{(2)} \leftarrow \underline{\underline{W}}^{(2)T}$$

$$\tilde{\underline{\underline{b}}}^{(1)} \leftarrow \underline{\underline{b}}^{(1)T}$$

$$\tilde{\underline{\underline{b}}}^{(2)} \leftarrow \underline{\underline{b}}^{(2)T}$$

Neural Networks for Classification

— How did we extend linear regression to logistic regression?

- By applying logistic (or sigmoid) function to the output of linear regression in case of binary classification
- For multi-class classification (with $y \in \{1, 2, \dots, M\}$ classes), we used the softmax function

$$\text{softmax}(\underline{z}) = \frac{1}{\sum_{j=1}^M e^{z_j}} \begin{bmatrix} e^{z_1} \\ e^{z_2} \\ \vdots \\ e^{z_M} \end{bmatrix}$$

— The softmax function now becomes an **additional** activation function, acting on the final layer of the neural network

$$\underline{h}^{(1)} = \sigma(\underline{w}^{(1)} \underline{x} + \underline{b}^{(1)})$$


$$\vdots$$
$$\underline{h}^{(L-1)} = \sigma(\underline{w}^{(L-1)} \underline{h}^{(L-2)} + \underline{b}^{(L-1)})$$

$$\underline{z}_{M \times 1} = \underline{w}^{(L)} \underline{h}^{(L-1)} + \underline{b}^{(L)}$$


$$\underline{g}_{M \times 1} = \text{softmax}(\underline{z})$$

MNIST example: Classification of handwritten digits

- Dataset has 60000 training images
- " " 10000 test/validation images
- Each data point consists of a 28×28 pixel grayscale image of a handwritten digit
- Each image is also labelled with the digit $0, 1, 2, \dots, 9$ that it depicts
- Each pixel intensity has been normalized between $[0, 1]$



- Consider image as input $\underline{x} = [x_1 \ x_2 \ \dots \ x_p]^T$
 - $p = 28 \times 28 = 784$ input variables (flattened out)
 - Each x_j corresponds to a pixel in the image and represents its intensity
 - $x_j = 0 \rightarrow$ black pixel
 - $x_j = 1 \rightarrow$ white pixel
 - Anything between 0 and 1 is grey pixel



- Consider image as input $\underline{x} = [x_1 \ x_2 \ \dots \ x_p]^T$
- $p = 28 \times 28 = 784$ input variables (flattened out)
- Each x_j corresponds to a pixel in the image and represents its intensity

$x_j = 0 \rightarrow$ black pixel

$x_j = 1 \rightarrow$ white pixel

Anything between 0 and 1 is grey pixel

Using a 2-layer NN


Consider 200 hidden units

$$\underline{H} = \sigma \left(\underbrace{\underline{X}}_{60000 \times 784} \underbrace{\underline{W}^{(1)T}}_{784 \times 200} + \underbrace{\underline{b}^{(1)T}}_{1 \times 200} \right)$$

60000×200

$$\underline{\hat{y}} = \underline{H} \underbrace{\underline{W}^{(2)T}}_{200 \times 10} + \underbrace{\underline{b}^{(2)T}}_{1 \times 10}$$

60000×10



$$\text{Total parameters} = 784 \times 200 + 200 + 200 \times 10 + 10 = 159010 !!$$